

Distributed Systems - Project Report

Diego Oniarti - Alessandra Dalla Verde

Year 2024-2025

Contents

1	Design Choices	1
1.1	System Layout	1
1.2	Messages	1
1.3	Round system	3
1.4	Crash and Recovery	3
2	Assumptions and Requirements	4
2.1	Sequential Consistency	4
2.2	Concurrent Transactions	4
3	Testing though Simulation	4

1 Design Choices

1.1 System Layout

The main classes of our system are the Nodes, the Clients, and the Coordinator.

The nodes and clients serve their roles as described in the project's assignment, while the coordinator is tasked to issue commands and enforce the project's assumptions.

To keep track of ongoing operations we implemented a "rounds system" in which the coordinator issues a group of operations, waits for their conclusion, and repeats the process.

1.2 Messages

The system's components communicate through messages, divided into two main categories: control messages and service messages.

Control messages are exchanged instantaneously between the coordinator and the other actors to issue commands (requests and management operations) and to get their results. With this information the coordinator can guarantee that nodes join and leave, crash and recover one at a time, and only when there are no ongoing operations.

Service messages implement the actual functionalities of the system. They simulate the messages that would be exchanged in a real world scenario, hence they are only sent between nodes and clients (not the coordinator). To make their behaviour more realistic they are sent with a random delay, mimicking the propagation delay of packets.¹

In **Fig.1** we illustrate all the messages of the system. To reduce the number of messages we adopted the following strategies:

¹The random delays *could* break FIFO, but we decided to ignore this as per the project's given assumptions

1. **Get** and **Set**: only the nodes responsible for the interested data item are contacted.
2. **Join**: the system only asks the joining node's clockwise neighbor for the data items that it is going to be responsible for, avoiding a number of unnecessary message exchanges. The latest version of these data items is then obtained through the Get operation. As stated above, reads already use the least amount of required messages.
3. **Leave**: the leaving node must transfer its data items to their new responsible nodes. If multiple data items are sent to the same node, they are bundled into the same message.
4. **Crash**: When a node crashes it sends no messages to the other nodes, only setting an internal flag. This operation is triggered by a control message.
5. **Recovery**: when a node enters recovery, it sends a data request to the other nodes to get updated on the status of the system. These messages are only sent to the $N - 1^2$ nodes in front and the $N - 1$ nodes behind the recovering one. Only these nodes may have the desired data items.

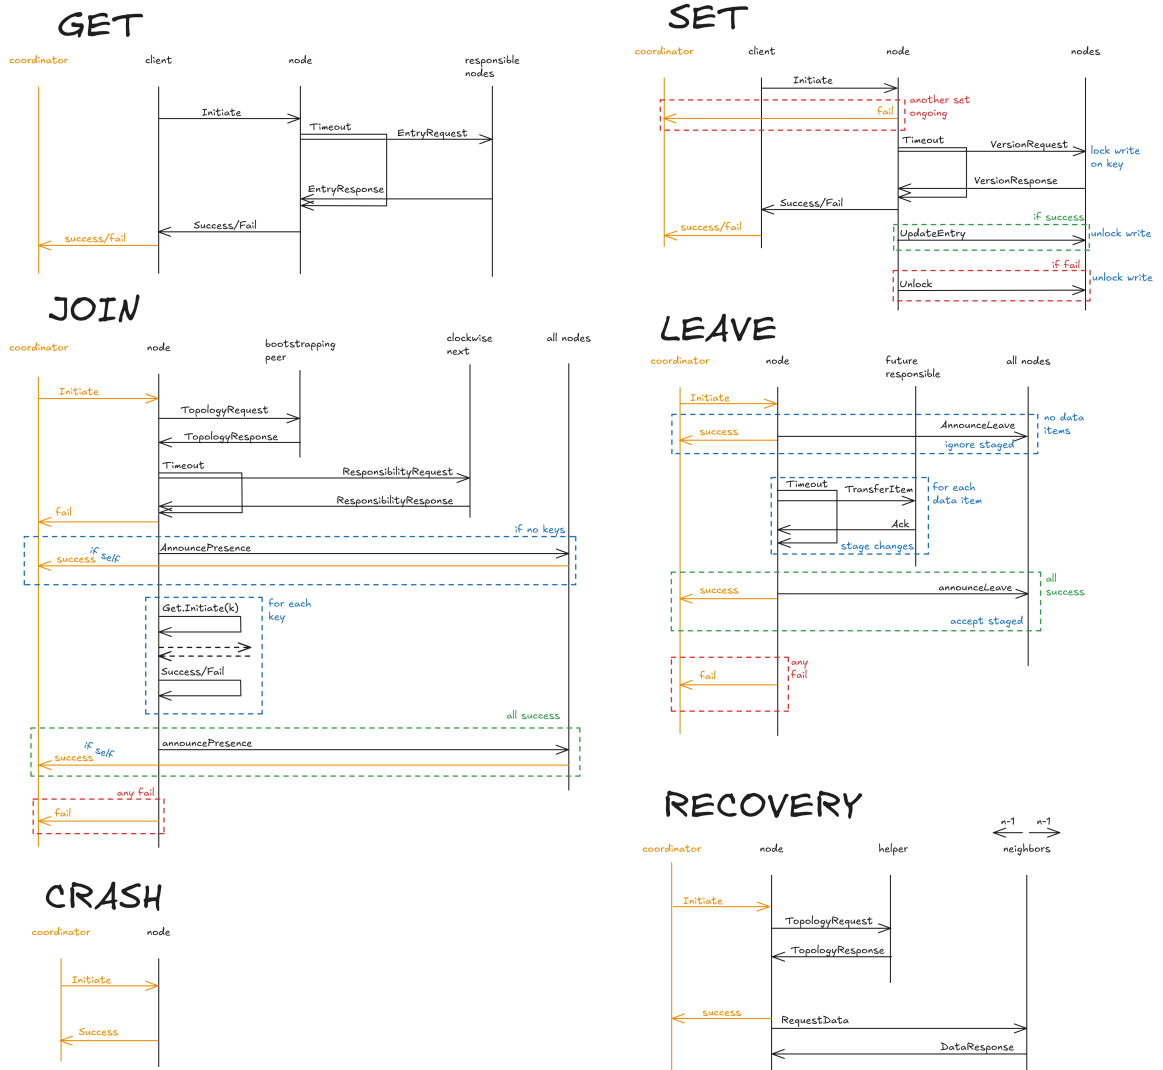


Figure 1: System messages

² N denotes the system's replication factor

1.3 Round system

There are two types of rounds: ones in which the system resolves a request (read or write) for each client, and ones in which a node performs management operations (join/leave, crash/recover).

Fig.2 shows the flowchart of a system's round. The first diagram represents the simulation's core: the coordinator continues issuing operations until it reaches a stopping criterion and decides which type of round it must begin. The second scheme is the rounds final part: the coordinator waits for the operations' results and an additional time delay to let the system settle. Only then it will start a new round.

Fundamental variables appear in the diagrams :

- **ROUNDS**: global variable indicating how many rounds the coordinator must perform
- **ongoing_actions**: number of pending operations for which the coordinator must receive a result (success or fail); it is set based on the type of round performed
- **nodes_in**: list of nodes in the network
- **nodes_out**: list of nodes out from the network from which the coordinator chooses a joining node and in which it puts leaving nodes
- **crashed_nodes**: list of crashed nodes

The coordinator updates the above lists based on the performed operations, thus at the begin of each round, the coordinator knows the nodes' state and it is able to perform actions respecting the system's requirements and assumptions.

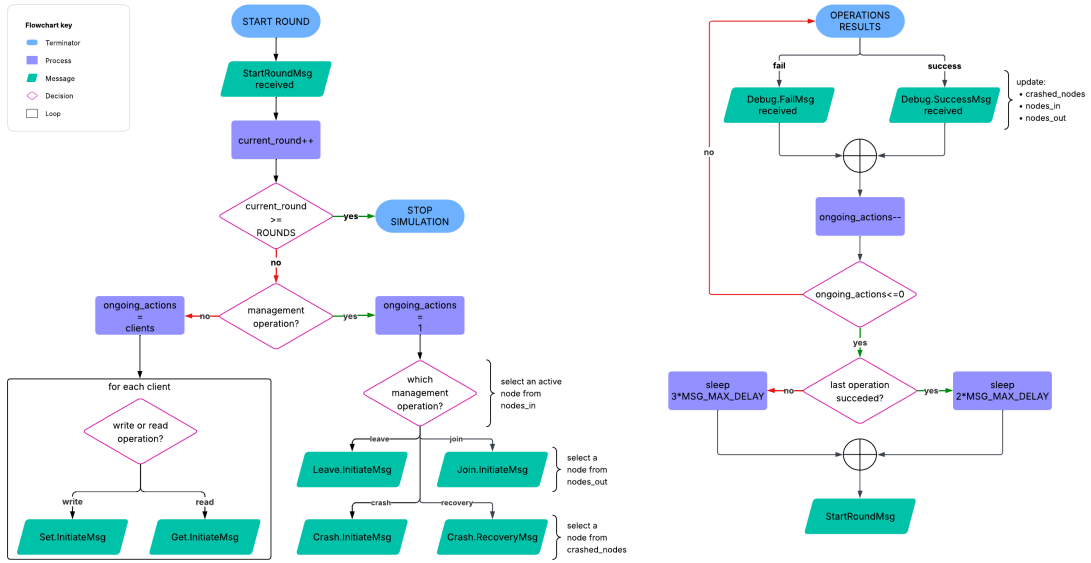


Figure 2: Round System

1.4 Crash and Recovery

When a node crashes, it simply set an internal flag (**crashed**) and it stops processing messages.

When a node recovers, it asks network topology and it requests updated versions of its data items. It may not receive one or more answers, in which case it will maintain its own versions. It is not a problem, because this would only happen if all other nodes responsible for a data item were also

crashed. If that were the case we can assume any other write operation would also have failed, without modifying the data items.

2 Assumptions and Requirements

To ensure the system's liveness, we made the following additional assumptions:

- clients contact an active node during set/get operations
- the joining node contacts an active bootstrapping peer
- join can fail when the node cannot read the data items it will be responsible for or when the clockwise neighbor is crashed; in this last case the system could stall
- leave can fail when the leaving node cannot share its data items within a given timeout
- a node can't leave if there are exactly N nodes in the network

2.1 Sequential Consistency

We provided sequential consistency by managing concurrent operations and by setting quorums appropriately.

Read/write and write/write conflicts (race conditions) are resolved by failing the read operation in the first case, and one of the two writes in the second one.

As we have seen in class, to provide sequential consistency with quorums, it is sufficient to set the parameters N , N_W , and N_R such that:

$$\begin{aligned} N_W &> \frac{N}{2} \\ N_W + N_R &> N \end{aligned}$$

2.2 Concurrent Transactions

Concurrent operations could affect the system's behavior. Write/write conflicts could cause inconsistency: responsible nodes could have the same version but different values for same data item. Read/write conflicts could break sequential consistency.

We manage these conflicts in the following way:

- write/write conflicts: the set transaction fails if the responsible nodes for it are managing other write operations. When a node receives a version request for a data item, it locks that item's key, and it no longer participates in any set operation on that item. The key is only freed when the related transaction ends (see messages of **set** operation in **Fig.1**). In some corner cases it may happen that all concurrent write operations fail. This behaviour occurs when none of the concurrent write operations manages to secure enough nodes for a quorum.
- read/write conflicts: the get transactions fail if the responsible nodes are participating in a set operation on the same data item. If the key is locked, the node does not continue performing the read operation.

3 Testing though Simulation

We perform two kinds of test on the system, both performed by the AppDebug class. Both tests rely on a log file containing a sequence of events in their global order.

The events stored in the file are:

- node storage modification: whenever data items are written or deleted from a node's storage, an entry describing the change is written to the file.
- operations: whenever an operation succeeds or fails, it is registered to the file as a single atomic entry.

First Test: Sequential Consistency A validator reads the log file and keeps track of all the successful read and write operations, storing each one as a node in a graph. The nodes are identified by the operation performed, the key of the item, its value, and its version.

Edges are then added between nodes to represent the happens-before relations between them:

1. write nodes are linked to all the read nodes with the same key and value, since a value has to be written before being read
2. if two read events for the same key happen on the same client, an edge is created to represent this ordering
3. each time a client writes something, an edge between all preceding read nodes (from the same client) and the newly formed write node is formed

Finally, the system checks for the existence of a topological order of this graph, and signals an error if it can't find any.

If none is found, it must be because the graph contains at least one loop, which shouldn't happen since the edges represent the happens-before relation which is acyclic.

If an ordering is found instead, the system enforced data-centric sequential consistency.

Second Test: Mirroring Execution For the second test, two replicas of the system are created, each one tracking the individual storages of all nodes.

The first replica perfectly mirrors the execution of the real system. Data items are added or removed from the nodes if and only if the log file states that an alteration happened to the node's storage.

The second replica only reads the entries of the log file that report a successful operation (be it a management one or a read/write request) and simulates it.

The simulation keeps track of which nodes are in the system, which ones are crashed, and updates their storages accordingly.

At the end of each round the storage of each node in the first replica is compared with the corresponding one in the second replica. If no divergence between the two systems is found we can confirm the actual system behaves correctly.